

Quick Notes - Feature Engineering and Cross Validation

Sometimes, model might fail to work as per expectation in the training data itself. A model may even appear to be performing very well in training but might fail to replicate the performance on unseen data, such as test data. What could be the reasons for such behaviour?

Sometimes the model might be good, but the data might not be perfect or might require some more pre-processing, or maybe the model performance isn't being tested the right way.

The video lectures are designed to help you understand techniques that will help deal with both the problems mentioned above using new pre-processing, oversampling, under sampling and regularization techniques.

The hands-on lectures discuss the implementation of the techniques using Python, and the primary focus will be on their implementation.